

# Exploratory Data Analysis

February 19, 2025

## 1 importing modules

```
[904]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from overview import load_bank_variables
from sklearn.preprocessing import LabelEncoder
from sklearn.preprocessing import RobustScaler

pd.set_option('display.max_colwidth', None) # Show full column content
pd.set_option('display.expand_frame_repr', False) # Disable line wrapping
pd.set_option('display.max_rows', None) # Show all rows
pd.set_option('display.max_columns', None) # Show all columns
pd.set_option('display.width', 1000) # Adjust column width
```

### 1.1 Loading Data

```
[905]: bank_df = pd.read_csv("bank.csv")
df = bank_df.copy()
```

```
[906]: load_bank_variables()
```

```
[906]: Variable Name
Description
0          age
Age
1          job  Type of job (e.g., 'admin.', 'blue-
collar', 'entrepreneur', 'housemaid', 'management', 'retired', 'self-
employed', 'services', 'student', 'technician', 'unemployed', 'unknown')
2          marital
Marital status (e.g., 'divorced', 'married', 'single', 'unknown'; 'divorced' means
divorced or widowed)
3          education  Education level (e.g., 'basic.4y
', 'basic.6y', 'basic.9y', 'high.school', 'illiterate', 'professional.course', 'univer
sity.degree', 'unknown')
4          default
```

Has credit in default?  
5 balance  
Average yearly balance (euros)  
6 housing  
Has housing loan?  
7 loan  
Has personal loan?  
8 contact  
Contact communication type (e.g., 'cellular','telephone')  
9 day\_of\_week  
Last contact day of the week  
10 month  
Last contact month of year (e.g., 'jan', 'feb', 'mar', ..., 'nov', 'dec')  
11 duration  
Last contact duration in seconds (only for benchmarks, discard for real prediction)  
12 campaign  
Number of contacts during this campaign (includes last contact)  
13 pdays  
Number of days since last contact from previous campaign (-1 means not contacted)  
14 previous  
Number of contacts before this campaign  
15 poutcome  
Outcome of previous campaign (e.g., 'failure','nonexistent','success')  
16 y  
Has the client subscribed to a term deposit?

## 1.2 Data Exploration

```
[907]: df.head(5)
```

```
[907]:
```

	age	job	marital	education	default	balance	housing	loan
contact	day	month	duration	campaign	pdays	previous	poutcome	subscribed
0	32.0	technician	single	tertiary	no	392	yes	no
cellular	1	apr	957	2	131	2	failure	no
1	39.0	technician	divorced	secondary	no	688	yes	yes
cellular	1	apr	233	2	133	1	failure	no
2	59.0	retired	married	secondary	no	1035	yes	yes
cellular	1	apr	126	2	239	1	failure	no
3	47.0	blue-collar	married	secondary	no	398	yes	yes
cellular	1	apr	274	1	238	2	failure	no
4	54.0	retired	married	secondary	no	1004	yes	no
cellular	1	apr	479	1	307	1	failure	no

```
[908]: df.tail(5)
```

```
[908]:      age      job marital education default balance housing loan
contact day month duration campaign pdays previous poutcome subscribed
1995  20.0      student   single      NaN      no      2785      no   no
cellular  16   sep      327      2      -1      0      NaN      yes
1996  28.0      admin.   single secondary      no      127      no   no
cellular  16   sep     1334      2      -1      0      NaN      yes
1997  81.0      retired married   primary      no     1154      no   no
telephone 17   sep      231      1      -1      0      NaN      yes
1998  46.0      services married   primary      no     4343     yes   no
NaN     20   sep      185      1      -1      0      NaN      yes
1999  40.0 entrepreneur married secondary      no     6403      no   no
cellular  22   sep      208      2      -1      0      NaN      yes
```

```
[909]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2000 entries, 0 to 1999
Data columns (total 17 columns):
#   Column      Non-Null Count  Dtype
---  -
0   age         1988 non-null   float64
1   job         1990 non-null   object
2   marital     2000 non-null   object
3   education   1896 non-null   object
4   default     2000 non-null   object
5   balance     2000 non-null   int64
6   housing     2000 non-null   object
7   loan        2000 non-null   object
8   contact     1809 non-null   object
9   day         2000 non-null   int64
10  month       2000 non-null   object
11  duration    2000 non-null   int64
12  campaign    2000 non-null   int64
13  pdays      2000 non-null   int64
14  previous    2000 non-null   int64
15  poutcome    1546 non-null   object
16  subscribed  2000 non-null   object
dtypes: float64(1), int64(6), object(10)
memory usage: 265.8+ KB
```

the dataset contains 2000 rows and 17 columns both numerical and categorical

- numerical : age , balance , duration , campaign , pdays , previous
- categorical: job , marital , education , default , housing , loan , contact , month , poutcome


there is some missing values in :

- age (12)
- job (10)
- education (104)

- contact (191)
- poutcome (454)

```
[910]: # Visualizing missing values
sns.heatmap(df.isnull(), cbar=False, cmap='viridis')
```

```
[910]: <Axes: >
```



### 1.3 seperate columns by type to plot each

```
[911]: numerical_columns = df.select_dtypes(include=['int64', 'float64']).columns
categorical_columns = df.select_dtypes(include=['object']).columns
```

### 1.4 stats of numerical column

```
[912]: df[numerical_columns].describe()
```

```
[912]:
```

	age	balance	day	duration	campaign
pdays	previous				
count	1988.000000	2000.000000	2000.000000	2000.000000	2000.000000
	2000.000000	2000.000000			
mean	41.753018	1413.663500	13.851500	292.020500	1.909500
	167.896000	2.561500			
std	12.724358	3131.224213	9.712189	221.557295	1.378862
	131.754126	3.400735			
min	18.000000	-980.000000	1.000000	7.000000	1.000000
	-1.000000	0.000000			
25%	32.000000	201.500000	5.000000	146.000000	1.000000
	75.750000	1.000000			
50%	38.000000	551.000000	12.000000	236.000000	1.000000
	182.000000	2.000000			
75%	50.000000	1644.500000	23.000000	379.000000	2.000000
	251.000000	3.000000			
max	93.000000	81204.000000	31.000000	1823.000000	11.000000
	854.000000	55.000000			


## 1.5 plotting numerical columns

```
[913]:
```

```
# Distribution Plot
fig, axes = plt.subplots(1, 2, figsize=(14, 5))
sns.histplot(df['balance'], kde=True, ax=axes[0])
axes[0].set_title(f"Distribution of balance")
axes[0].set_xlabel(column)
axes[0].set_ylabel("Frequency")
axes[0].set_xlim(left=-10000)

# Boxplot
sns.boxplot(x=df["balance"], ax=axes[1])
axes[1].set_title(f"Boxplot of {column}")

plt.tight_layout()
plt.show()
```

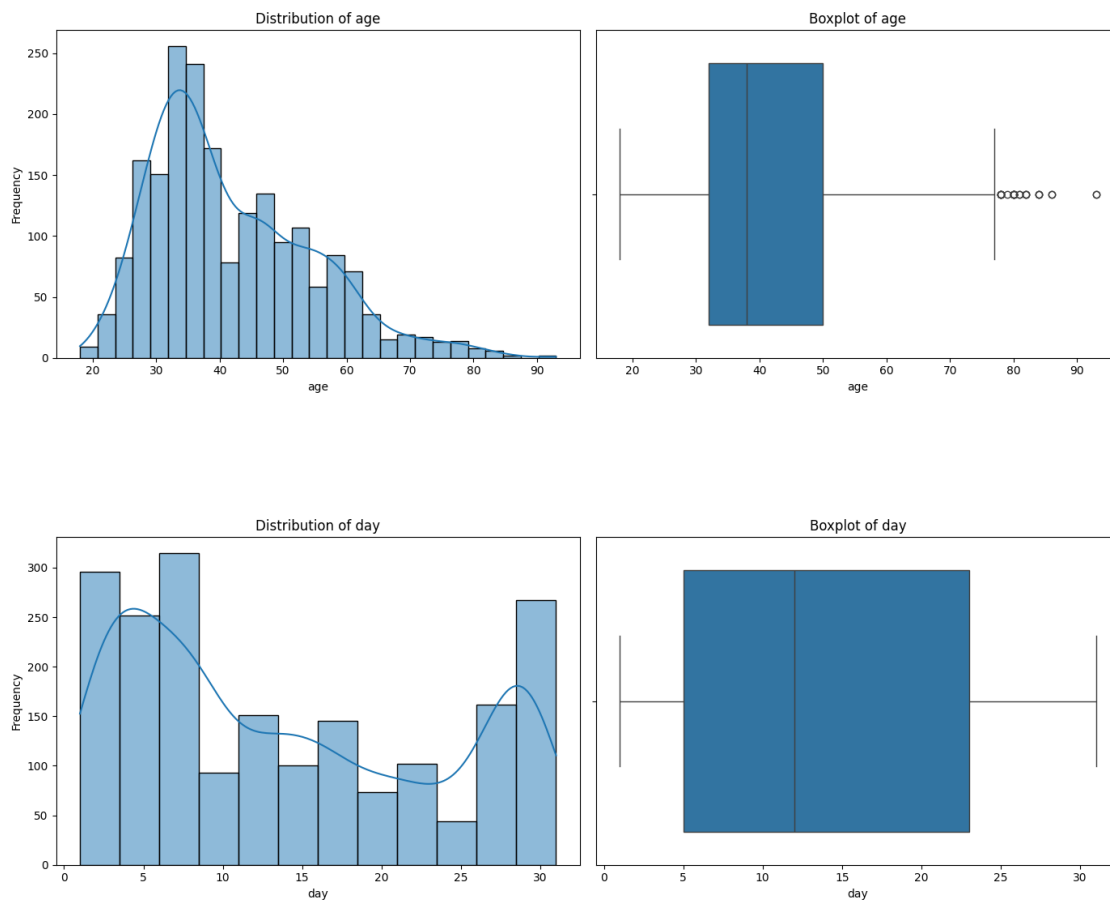



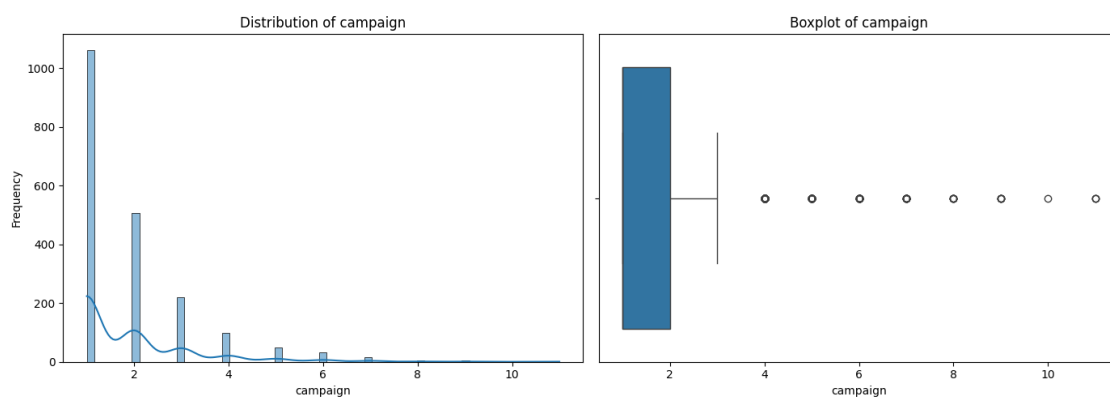

```
[914]: for column in numerical_columns:
    if column != 'balance': # Skip the 'balance' column
        fig, axes = plt.subplots(1, 2, figsize=(14, 5))


        # Distribution Plot
        sns.histplot(df[column].dropna(), kde=True, ax=axes[0])
        axes[0].set_title(f"Distribution of {column}")
        axes[0].set_xlabel(column)
        axes[0].set_ylabel("Frequency")

        # Boxplot
        sns.boxplot(x=df[column], ax=axes[1])
        axes[1].set_title(f"Boxplot of {column}")

    plt.tight_layout()
    plt.show()
```







```
[ ]:
```

## 1.6 unique values of categorical variables

```
[915]: for column in df[categorical_columns]:
        print(f"{column} :")
        print(df[column].unique(), "\n")
```

```
job :
['technician' 'retired' 'blue-collar' 'self-employed' 'services'
 'management' 'admin.' 'unemployed' 'student' 'entrepreneur' 'housemaid'
 nan]
```

```
marital :
['single' 'divorced' 'married']
```

```
education :
['tertiary' 'secondary' nan 'primary']
```

```
default :
['no' 'yes']
```

```
housing :
['yes' 'no']
```

```
loan :
['no' 'yes']
```

```
contact :
['cellular' 'telephone' nan]
```

```
month :
['apr' 'dec' 'feb' 'jan' 'mar' 'may' 'nov' 'oct' 'aug' 'jul' 'jun' 'sep']
```



```
poutcome :
['failure' 'other' 'success' nan]
```

```
subscribed :
['no' 'yes']
```

nan (missing value) exists on these features: \* Job \* Education \* Contact \* Poutcome

```
[916]: df[categorical_columns].describe()
```

```
[916]:
```

	job	marital	education	default	housing	loan	contact	month
poutcome	subscribed							
count	1990	2000	1896	2000	2000	2000	1809	2000
1546	2000							
unique	11	3	3	2	2	2	2	12
3	2							
top	management	married	secondary	no	no	no	cellular	feb
failure	no							
freq	461	1111	995	1985	1037	1750	1663	404
955	1000							





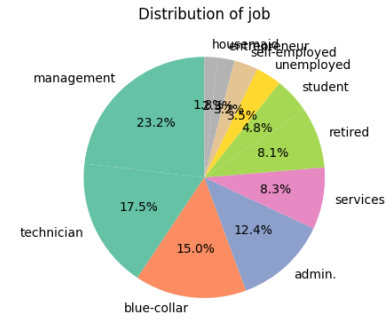

## 1.7 plotting categorical columns


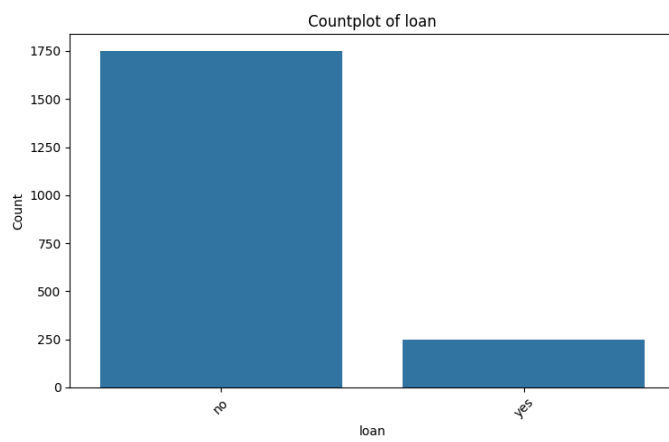




```
[917]: for column in categorical_columns:
        fig, axes = plt.subplots(1, 2, figsize=(14, 5)) # 1 row, 2 columns


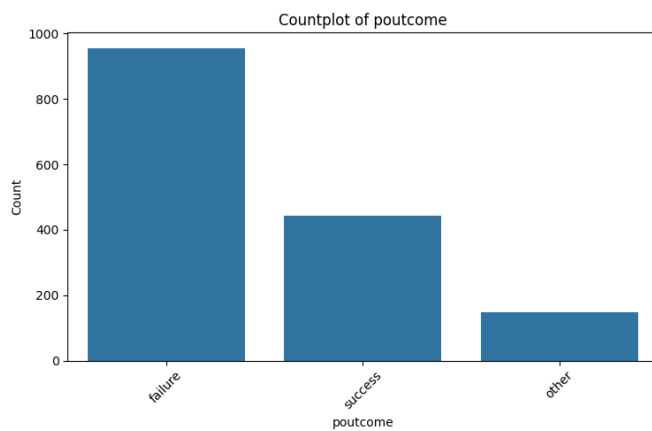

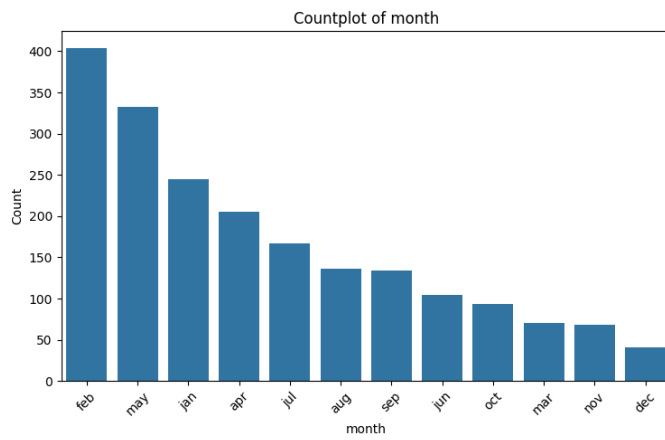

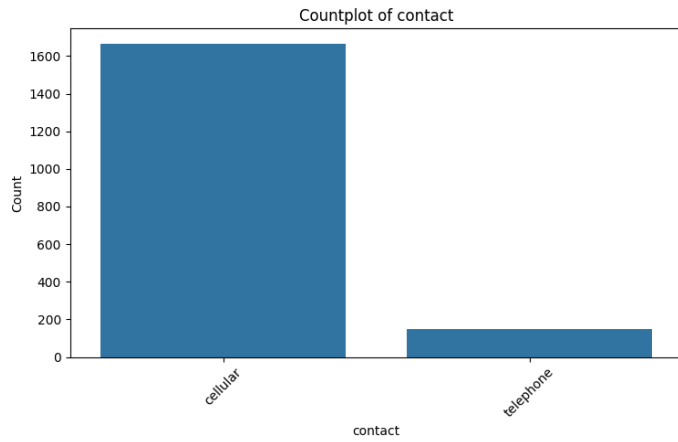
        # Countplot
        sns.countplot(data=df, x=column, order=df[column].value_counts().index,
        ↪ax=axes[0])
        axes[0].set_title(f"Countplot of {column}")
        axes[0].set_xlabel(column)
        axes[0].set_ylabel("Count")
        axes[0].tick_params(axis='x', rotation=45)


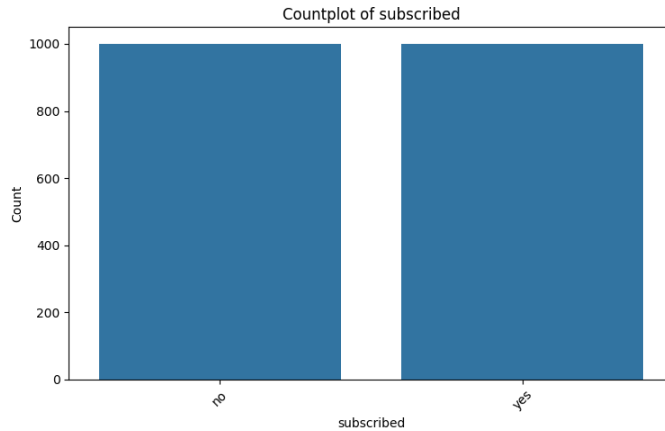
        # Pie Chart
        df[column].value_counts().plot.pie(autopct='%1.1f%%', ax=axes[1],
        ↪startangle=90, cmap='Set2')
        axes[1].set_title(f"Distribution of {column}")
        axes[1].set_ylabel('')

        plt.tight_layout()
        plt.show()
```









contact , default show low variaty loan show a little variaty but still might contribute to the data


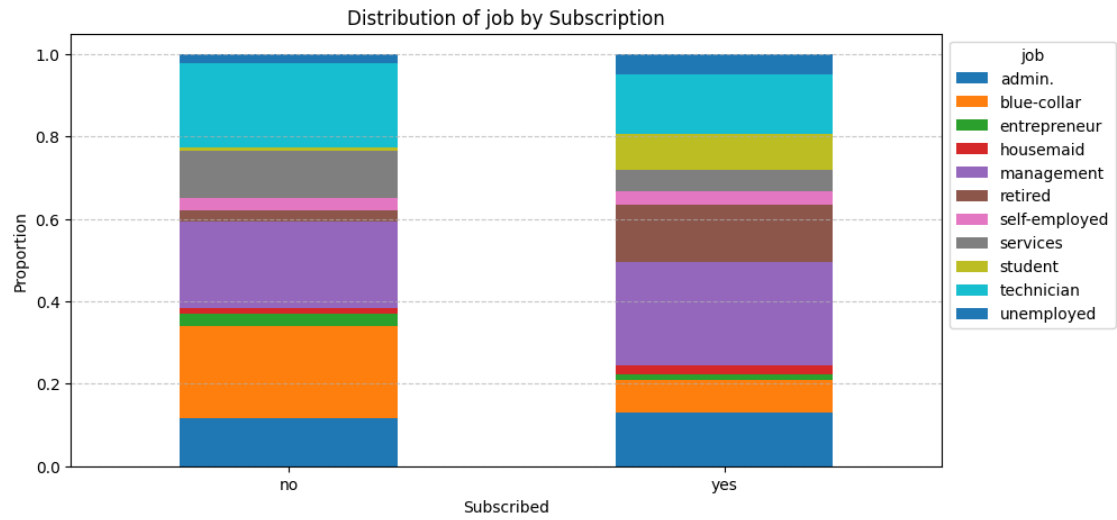
## 1.8 targeted comparison

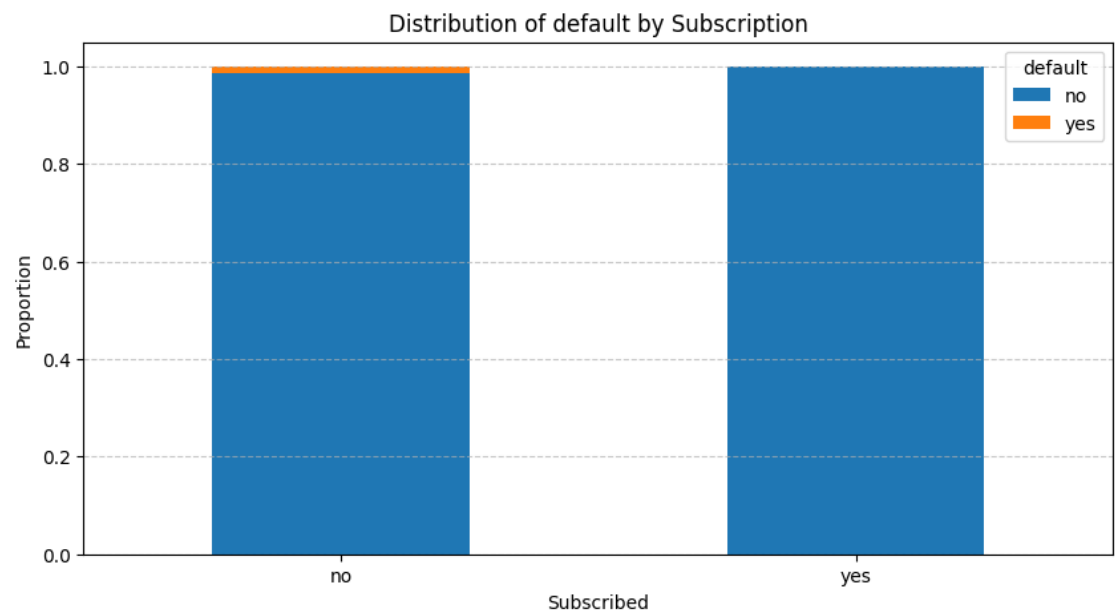
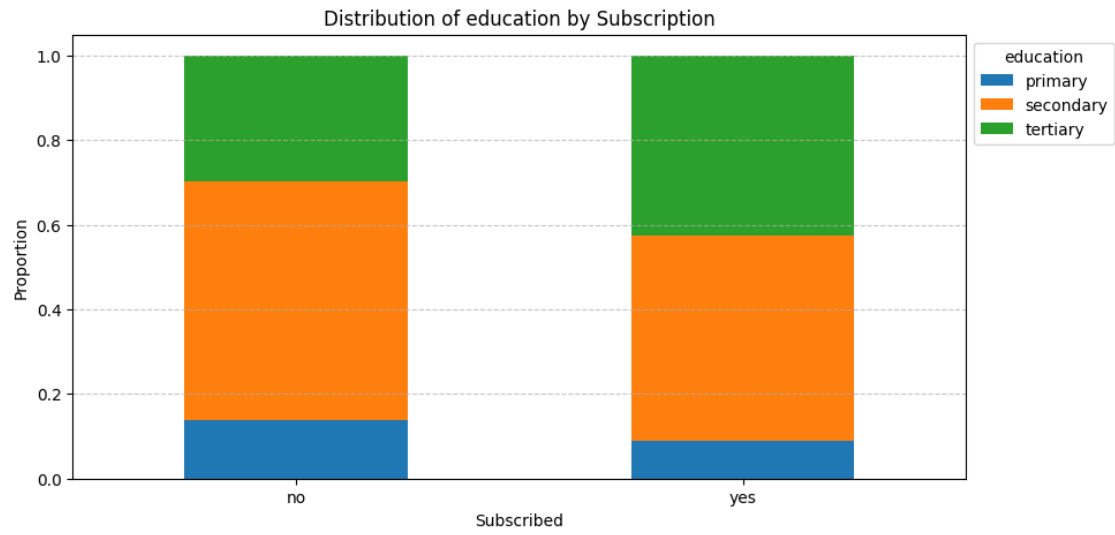
```
[918]: # Group by 'subscribed' and plot value counts for all categorical columns
        ↳ together
for column in categorical_columns:
    if column != 'subscribed':
        grouped = df.groupby('subscribed')[column].value_counts(normalize=True).
        ↳ unstack().fillna(0)

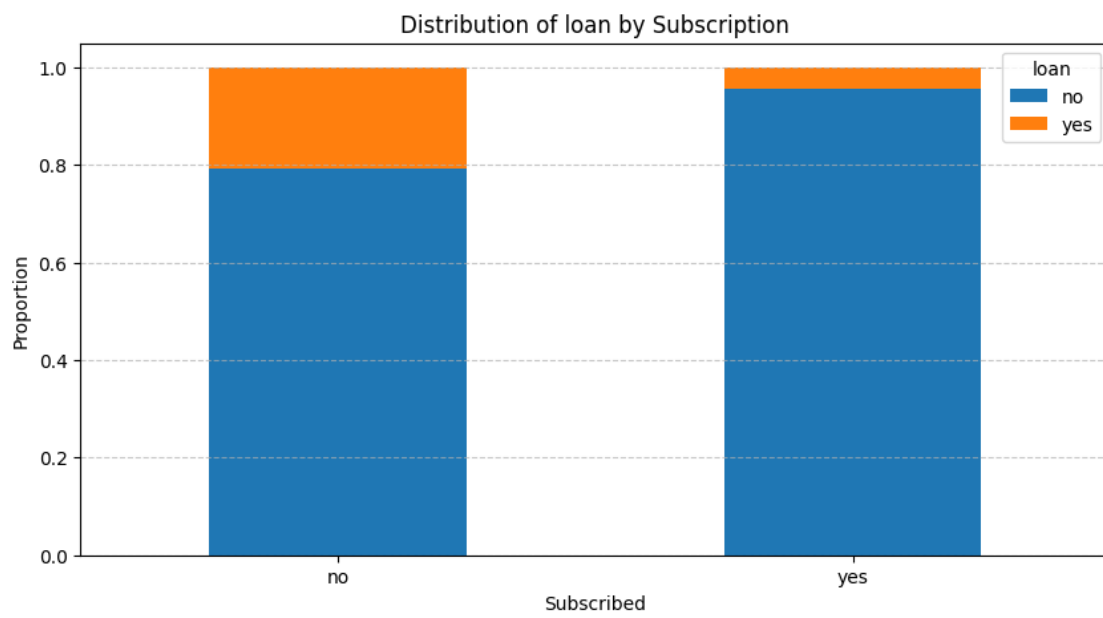
        # Plot as a stacked bar chart
        grouped.plot(kind='bar', stacked=True, figsize=(10, 5))

        plt.title(f"Distribution of {column} by Subscription")
        plt.xlabel("Subscribed")
        plt.ylabel("Proportion")
        plt.legend(title=column, bbox_to_anchor=(1, 1))
        plt.xticks(rotation=0)
        plt.grid(axis="y", linestyle="--", alpha=0.7)



        plt.show()
```














```
[919]: print("--- Numerical Columns Grouped by Subscribed ---")

for column in numerical_columns:
    print(f"\nStatistics for {column}:\n")
    display(df.groupby('subscribed')[column].describe())
```

--- Numerical Columns Grouped by Subscribed ---

Statistics for age:

	count	mean	std	min	25%	50%	75%	max
subscribed								
no	991.0	40.655903	9.192425	22.0	33.0	39.0	48.0	64.0
yes	997.0	42.843531	15.382656	18.0	31.0	38.0	54.0	93.0

Statistics for balance:

	count	mean	std	min	25%	50%	75%	max
subscribed								
no	1000.0	942.862	2007.134003	-980.0	114.75	393.0	970.25	26306.0
yes	1000.0	1884.465	3891.864047	-205.0	315.00	875.0	2304.50	81204.0

Statistics for day:

	count	mean	std	min	25%	50%	75%	max
subscribed								

no	1000.0	12.364	10.667394	1.0	4.0	8.0	27.25	30.0
yes	1000.0	15.339	8.397893	1.0	9.0	14.0	22.00	31.0

Statistics for duration:

	count	mean	std	min	25%	50%	75%	max
subscribed								
no	1000.0	206.696	175.152259	7.0	96.0	155.5	256.00	1823.0
yes	1000.0	377.345	230.154246	23.0	224.0	310.0	457.25	1720.0

Statistics for campaign:

	count	mean	std	min	25%	50%	75%	max
subscribed								
no	1000.0	1.957	1.443341	1.0	1.0	1.0	2.0	11.0
yes	1000.0	1.862	1.310219	1.0	1.0	1.0	2.0	11.0

Statistics for pdays:

	count	mean	std	min	25%	50%	75%	max
subscribed								
no	1000.0	185.400	99.759611	-1.0	136.0	211.0	259.0	536.0
yes	1000.0	150.392	155.468012	-1.0	-1.0	123.5	185.0	854.0

Statistics for previous:

	count	mean	std	min	25%	50%	75%	max
subscribed								
no	1000.0	2.362	3.287516	0.0	1.0	2.0	3.0	51.0
yes	1000.0	2.761	3.500590	0.0	0.0	2.0	4.0	55.0

## 1.9 targetting the pdays

```
[920]: (df['pdays'] == -1).mean()
```

```
[920]: np.float64(0.227)
```

22.7 % of the clients hasn t been contacted

## 1.10 targeting the balance

```
[921]: # How many zero values in `balance`?
print("There are %d account holder or %5f of the total clients who have zero_
↳balance" % ((df[df['balance']==0]['balance'].count()),
↳
↳      (df[df['balance']==0]['balance'].count()/(df['balance'].count()))
# How many negative values in `balance`?
print("There are %d account holder or %5f of the total clients who owe money" %_
↳((df[df['balance']<0]['balance'].count()),
↳
↳      (df[df['balance']<0]['balance'].count()/(df['balance'].count()))
```

There are 86 account holder or 0.043000 of the total clients who have zero balance

There are 93 account holder or 0.046500 of the total clients who owe money

```
[922]: # Is there any of those clients who subscribed to term deposit?
print("There are %d account holder who have zero balance and subscribed term_
↳deposit" % df[(df['balance']==0) & (df['subscribed']=='yes')]['balance'].
↳count())

print("There are %d account holder who have negative balance and subscribed_
↳term deposit" % df[(df['balance']<0) & (df['subscribed']=='yes')]['balance'].
↳count())
```

There are 42 account holder who have zero balance and subscribed term deposit

There are 7 account holder who have negative balance and subscribed term deposit

## 1.11 targeting duration

```
[923]: # The range of calls duration from clients who subscribed to the term deposit
print("The minimum duration (in seconds) to finalize a deal :",_
↳df[df['subscribed']=='yes']['duration'].min())
print("The maximum duration (in seconds) to finalize a deal :",_
↳df[df['subscribed']=='yes']['duration'].max())

# The average of calls duration from clients who subscribed to the term deposit
print("The average duration (in seconds) to finalize a deal :",_
↳df[df['subscribed']=='yes']['duration'].mean())
```


The minimum duration (in seconds) to finalize a deal : 23

The maximum duration (in seconds) to finalize a deal : 1720

The average duration (in seconds) to finalize a deal : 377.345

```
[924]: # The calls duration distribution from clients who subscribed to the term_
↳deposit
sns.histplot(df[df['subscribed']=='yes']['duration'], color='blue', kde=True)
```

[924]: <Axes: xlabel='duration', ylabel='Count'>



Most subscribing clients had calls lasting between 175 and 375 seconds.

!!!!!!!!! correlation doesn't mean causation

```
[925]: # show the clients with no contact in the last campaign but subscribed to the term deposit
df[(df['duration']==0) & (df['subscribed']=='yes')]
```

```
[925]: Empty DataFrame
Columns: [age, job, marital, education, default, balance, housing, loan, contact, day, month, duration, campaign, pdays, previous, poutcome, subscribed]
Index: []
```

## 1.12 targeting the poutcome

```
[926]: # Show Clients who have been contacted before but unknown outcome
df[(df['poutcome'] == '') & (df['previous']!=0)]
```

```
[926]: Empty DataFrame
Columns: [age, job, marital, education, default, balance, housing, loan,
```

```
contact, day, month, duration, campaign, pdays, previous, poutcome, subscribed]
Index: []
```

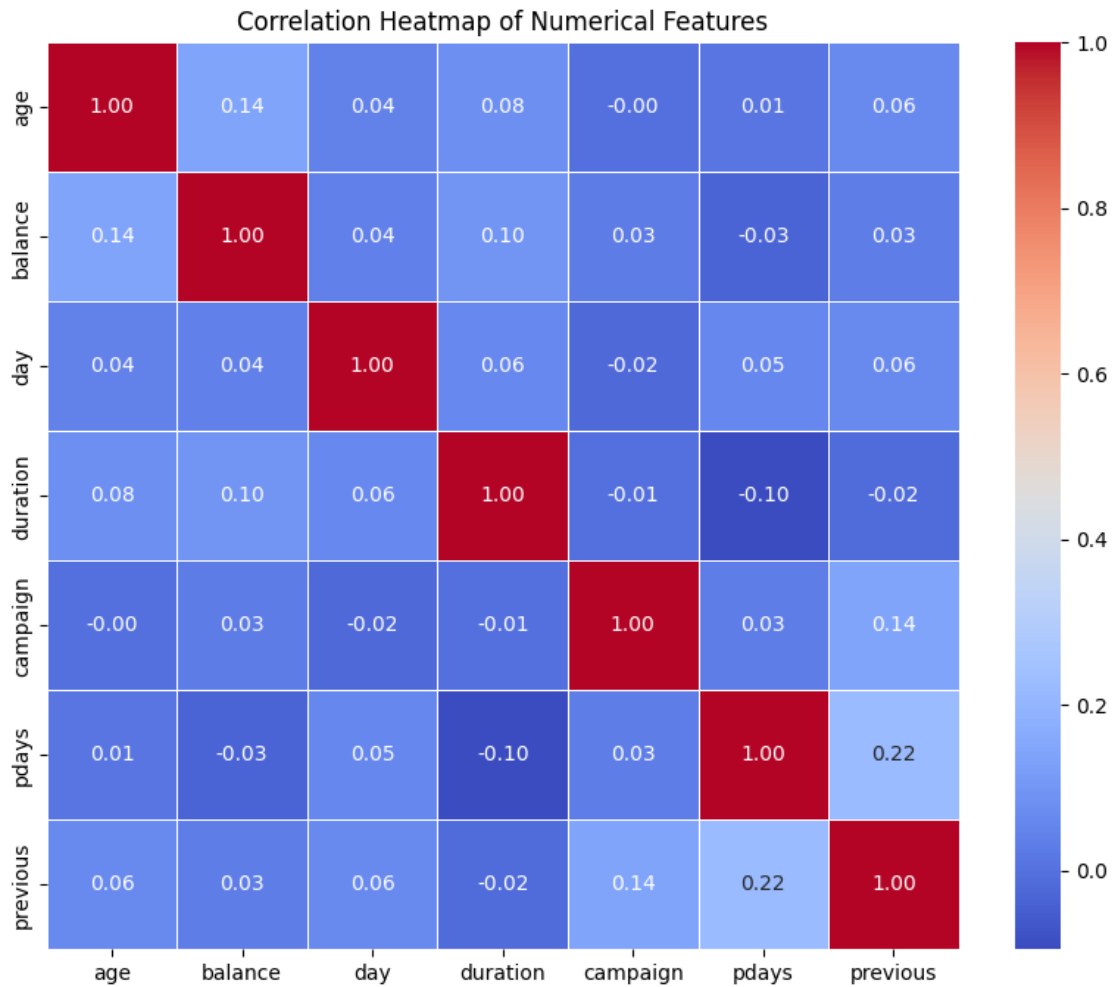
```
[927]: # pdays is -1 => previous is 0 ?
print(df[(df['pdays'] != -1) & (df['previous']==0)]['pdays'].count() ==
      df[(df['pdays'] == -1) & (df['previous'] != 0)]['pdays'].count())
```

True

### 1.13 correlation matrix

```
[928]: # Calculate the correlation matrix
correlation_matrix = df.select_dtypes(include=['int64', 'float64']).corr()

# Plot the heatmap
plt.figure(figsize=(10, 8))
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm', fmt='.2f',
            linewidths=0.5)
plt.title('Correlation Heatmap of Numerical Features')
plt.show()
```



there is almost no correlation between any pair of numerical columns

## 2 summary

nan (missing value) exists on these features: \* Job \* Education \* Contact \* Poutcome

columns to drop : \* contact : Almost all clients contacted via cellular

\* default : Almost no clients have defaulted (99.2% “no”)

\* day : we ll treat the cycle data in month and dispose of this column \* duration : has a wide range of values ,might outweigh other important features.

## 3 DATA preparation

```
[929]: df.isnull().sum()
```

```
[929]: age          12
      job          10
      marital       0
      education    104
      default       0
      balance       0
      housing       0
      loan          0
      contact     191
      day           0
      month         0
      duration      0
      campaign      0
      pdays        0
      previous      0
      poutcome     454
      subscribed    0
      dtype: int64
```

```
[930]: original_df = df.copy()
```

### 3.0.1 Dropping unnecessary columns

```
[931]: df.drop(columns=['contact', 'duration', 'default', 'day'], inplace=True)
```

```
[932]: df.columns
```

```
[932]: Index(['age', 'job', 'marital', 'education', 'balance', 'housing', 'loan',
      'month', 'campaign', 'pdays', 'previous', 'poutcome', 'subscribed'],
      dtype='object')
```

### 3.0.2 handling missing values

- age : the distribution is positively skewed for that we ll use the median to impute the missing values

```
[933]: df.fillna({'age': df['age'].median()}, inplace=True)
```

- poutcome : the data isn't missing but rather unknown

```
[934]: df['poutcome'] = df['poutcome'].replace("", "never")
```

- job : missing values < 0.5 % of the data , we can dispose of those rows

```
[935]: df = df.dropna(subset=['job'])
```

- education : we create a new category of unknown education level ' unknown '

```
[936]: df.fillna({'education': 'unknown'}, inplace=True)
```



### 3.0.3 encoding

```
[937]: # Convert 'month' into seasonal categories
season_map = {
    'dec': 'Winter', 'jan': 'Winter', 'feb': 'Winter',
    'mar': 'Spring', 'apr': 'Spring', 'may': 'Spring',
    'jun': 'Summer', 'jul': 'Summer', 'aug': 'Summer',
    'sep': 'Fall', 'oct': 'Fall', 'nov': 'Fall'
}

df['season'] = df['month'].map(season_map)
df.drop(columns=['month'], inplace=True) # Drop original month column

# One-Hot Encoding for 'season' column
df = pd.get_dummies(df, columns=['season'], drop_first=True)

# One-Hot Encoding for Other Nominal Categorical Columns
df = pd.get_dummies(df, columns=['job', 'marital', 'education', 'poutcome'],
                    drop_first=True)

[938]: # Label Encoding for Binary Categorical Columns
binary_cols = ['housing', 'loan']
le = LabelEncoder()

for col in binary_cols:
    df[col] = le.fit_transform(df[col])
```

### 3.1 Scaling

```
[ ]: # Define numerical columns to scale
numerical_columns = ['age', 'balance', 'campaign', 'previous', 'pdays']

# Initialize RobustScaler
scaler = RobustScaler()

# Apply scaling
df[numerical_columns] = scaler.fit_transform(df[numerical_columns])
```