

Exploratory Data Analysis

February 19, 2025

1 importing modules

```
[1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from overview import load_bank_variables
from sklearn.preprocessing import LabelEncoder
from sklearn.preprocessing import RobustScaler
from sklearn.decomposition import PCA
from sklearn.cluster import KMeans

pd.set_option('display.max_colwidth', None) # Show full column content
pd.set_option('display.expand_frame_repr', False) # Disable line wrapping
pd.set_option('display.max_rows', None) # Show all rows
pd.set_option('display.max_columns', None) # Show all columns
pd.set_option('display.width', 1000) # Adjust column width
```

1.1 Loading Data

```
[2]: bank_df = pd.read_csv("bank.csv")
df = bank_df.copy()
```

1.2 variables

```
[3]: load_bank_variables()
```

```
[3]: Variable Name
Description
0          age
Age
1          job  Type of job (e.g., 'admin.', 'blue-
collar', 'entrepreneur', 'housemaid', 'management', 'retired', 'self-
employed', 'services', 'student', 'technician', 'unemployed', 'unknown')
2          marital
```


| | | | | | | | | |
|----------|------|---------|---------|-----------|-----|------|---------|----|
| cellular | 1 | apr | 274 | 1 | 238 | 2 | failure | no |
| 4 | 54.0 | retired | married | secondary | no | 1004 | yes | no |
| cellular | 1 | apr | 479 | 1 | 307 | 1 | failure | no |

```
[5]: df.tail(5)
```

```
[5]:
```

| | age | job | marital | education | default | balance | housing | loan | |
|-----------|------|-------|--------------|-----------|-----------|----------|----------|------------|----|
| contact | day | month | duration | campaign | pdays | previous | poutcome | subscribed | |
| 1995 | 20.0 | | student | single | NaN | no | 2785 | no | no |
| cellular | 16 | sep | 327 | 2 | -1 | 0 | NaN | yes | |
| 1996 | 28.0 | | admin. | single | secondary | no | 127 | no | no |
| cellular | 16 | sep | 1334 | 2 | -1 | 0 | NaN | yes | |
| 1997 | 81.0 | | retired | married | primary | no | 1154 | no | no |
| telephone | 17 | sep | 231 | 1 | -1 | 0 | NaN | yes | |
| 1998 | 46.0 | | services | married | primary | no | 4343 | yes | no |
| NaN | 20 | sep | 185 | 1 | -1 | 0 | NaN | yes | |
| 1999 | 40.0 | | entrepreneur | married | secondary | no | 6403 | no | no |
| cellular | 22 | sep | 208 | 2 | -1 | 0 | NaN | yes | |

```
[6]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2000 entries, 0 to 1999
Data columns (total 17 columns):
#   Column          Non-Null Count  Dtype
---  -
0   age             1988 non-null   float64
1   job             1990 non-null   object
2   marital         2000 non-null   object
3   education       1896 non-null   object
4   default         2000 non-null   object
5   balance         2000 non-null   int64
6   housing         2000 non-null   object
7   loan           2000 non-null   object
8   contact         1809 non-null   object
9   day             2000 non-null   int64
10  month           2000 non-null   object
11  duration        2000 non-null   int64
12  campaign        2000 non-null   int64
13  pdays           2000 non-null   int64
14  previous        2000 non-null   int64
15  poutcome       1546 non-null   object
16  subscribed      2000 non-null   object
dtypes: float64(1), int64(6), object(10)
memory usage: 265.8+ KB
```

the dataset contains 2000 rows and 17 columns both numerical and categorical

- numerical : age , balance , duration , campaign , pdays , previous


- categorical: job , marital , education , default , housing , loan , contact , month , poutcome

there is some missing values in :

- age (12)
- job (10)
- education (104)
- contact (191)
- poutcome (454)

```
[7]: # Visualizing missing values
sns.heatmap(df.isnull(), cbar=False, cmap='viridis')
```

[7]: <Axes: >



1.4 separate columns by type to plot each

```
[8]: numerical_columns = df.select_dtypes(include=['int64', 'float64']).columns
categorical_columns = df.select_dtypes(include=['object']).columns
```

1.5 stats of numerical column

```
[9]: df[numerical_columns].describe()
```

```
[9]:
```


| | age | balance | day | duration | campaign |
|-------|-------------|--------------|-------------|-------------|-------------|
| pdays | previous | | | | |
| count | 1988.000000 | 2000.000000 | 2000.000000 | 2000.000000 | 2000.000000 |
| | 2000.000000 | 2000.000000 | | | |
| mean | 41.753018 | 1413.663500 | 13.851500 | 292.020500 | 1.909500 |
| | 167.896000 | 2.561500 | | | |
| std | 12.724358 | 3131.224213 | 9.712189 | 221.557295 | 1.378862 |
| | 131.754126 | 3.400735 | | | |
| min | 18.000000 | -980.000000 | 1.000000 | 7.000000 | 1.000000 |
| | -1.000000 | 0.000000 | | | |
| 25% | 32.000000 | 201.500000 | 5.000000 | 146.000000 | 1.000000 |
| | 75.750000 | 1.000000 | | | |
| 50% | 38.000000 | 551.000000 | 12.000000 | 236.000000 | 1.000000 |
| | 182.000000 | 2.000000 | | | |
| 75% | 50.000000 | 1644.500000 | 23.000000 | 379.000000 | 2.000000 |
| | 251.000000 | 3.000000 | | | |
| max | 93.000000 | 81204.000000 | 31.000000 | 1823.000000 | 11.000000 |
| | 854.000000 | 55.000000 | | | |

1.6 plotting numerical columns

```
[10]: # Distribution Plot
fig, axes = plt.subplots(1, 2, figsize=(14, 5))
sns.histplot(df['balance'], kde=True, ax=axes[0])
axes[0].set_title(f"Distribution of balance")
axes[0].set_xlabel("balance")
axes[0].set_ylabel("Frequency")
axes[0].set_xlim(left=-10000)

# Boxplot
sns.boxplot(x=df["balance"], ax=axes[1])
axes[1].set_title(f"Boxplot of balance")

plt.tight_layout()
plt.show()
```






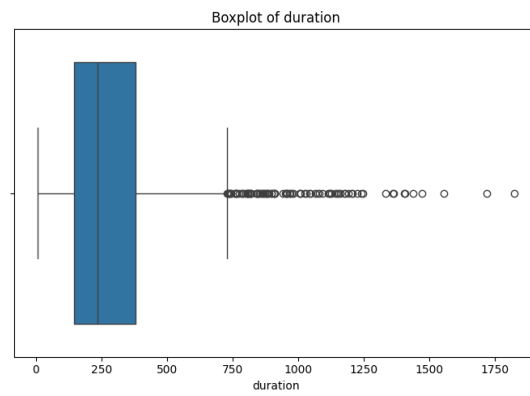



```
[11]: for column in numerical_columns:
    if column != 'balance': # Skip the 'balance' column
        fig, axes = plt.subplots(1, 2, figsize=(14, 5))

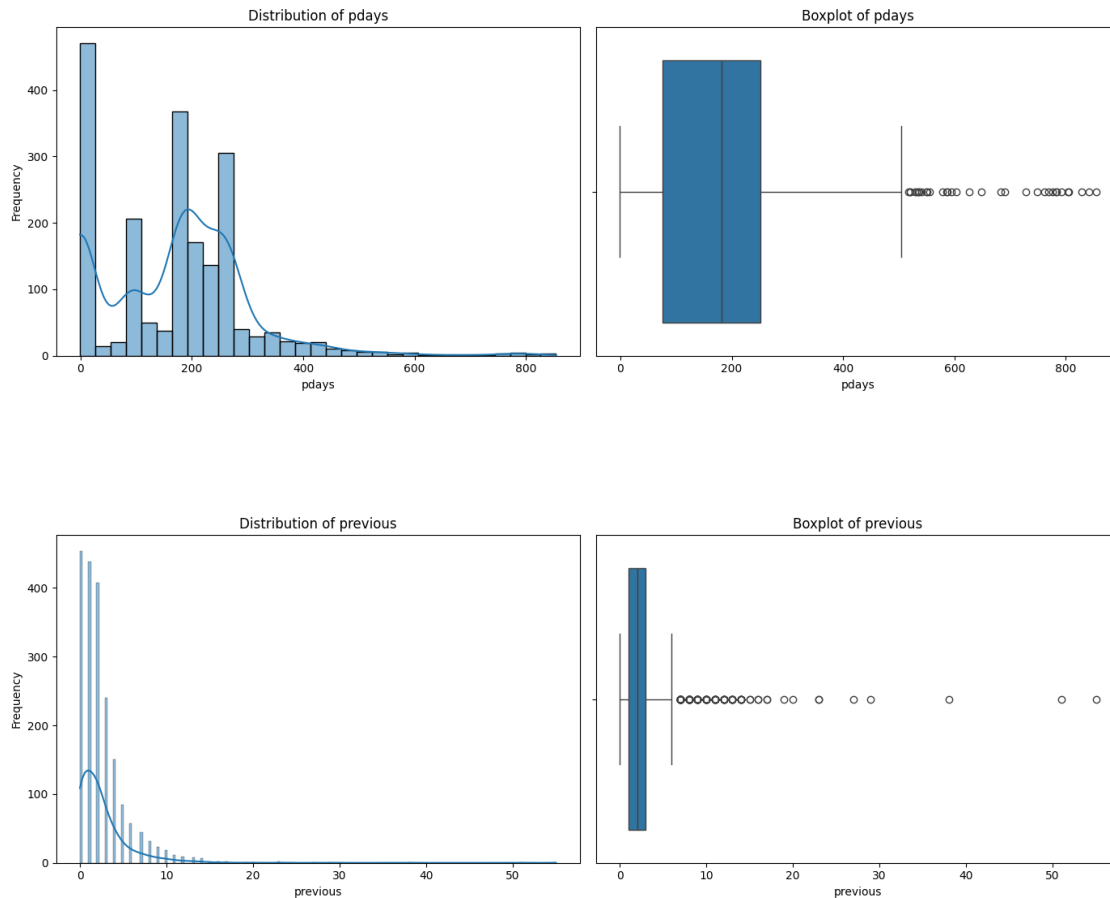
        # Distribution Plot
        sns.histplot(df[column].dropna(), kde=True, ax=axes[0])
        axes[0].set_title(f"Distribution of {column}")
        axes[0].set_xlabel(column)
        axes[0].set_ylabel("Frequency")

        # Boxplot
        sns.boxplot(x=df[column], ax=axes[1])
        axes[1].set_title(f"Boxplot of {column}")

    plt.tight_layout()
    plt.show()
```







[]:

1.7 unique values of categorical variables

```
[12]: for column in df[categorical_columns]:
       print(f"{column} :")
       print(df[column].unique(), "\n")
```

```
job :
['technician' 'retired' 'blue-collar' 'self-employed' 'services'
 'management' 'admin.' 'unemployed' 'student' 'entrepreneur' 'housemaid'
 nan]
```

```
marital :
['single' 'divorced' 'married']
```

```
education :
['tertiary' 'secondary' nan 'primary']
```



```

default :
['no' 'yes']

housing :
['yes' 'no']

loan :
['no' 'yes']

contact :
['cellular' 'telephone' nan]

month :
['apr' 'dec' 'feb' 'jan' 'mar' 'may' 'nov' 'oct' 'aug' 'jul' 'jun' 'sep']

poutcome :
['failure' 'other' 'success' nan]

subscribed :
['no' 'yes']

```

nan (missing value) exists on these features: * Job * Education * Contact * Poutcome

```
[13]: df[categorical_columns].describe()
```

```

[13]:          job  marital  education default housing  loan  contact month
poutcome subscribed
count          1990      2000          1896      2000      2000      2000          1809      2000
1546          2000
unique           11         3           3         2         2         2           2         12
3              2
top    management  married  secondary      no      no      no  cellular    feb
failure          no
freq           461      1111          995      1985      1037      1750          1663      404
955          1000

```

1.8 plotting categorical columns

```

[14]: for column in categorical_columns:
        fig, axes = plt.subplots(1, 2, figsize=(14, 5)) # 1 row, 2 columns

        # Countplot
        sns.countplot(data=df, x=column, order=df[column].value_counts().index,
            ↪ax=axes[0])
        axes[0].set_title(f"Countplot of {column}")
        axes[0].set_xlabel(column)
        axes[0].set_ylabel("Count")

```

```
axes[0].tick_params(axis='x', rotation=45)
```

```
# Pie Chart
```


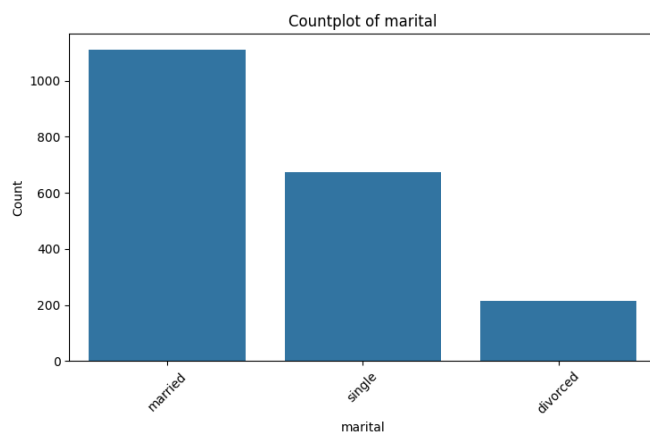
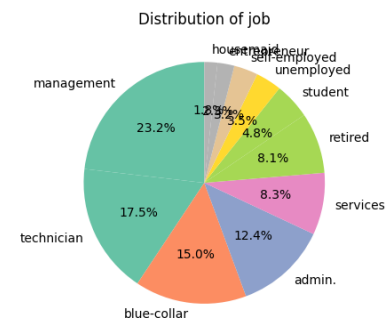

```
df[column].value_counts().plot.pie(autopct='%1.1f%%', ax=axes[1],
↪startangle=90, cmap='Set2')
```




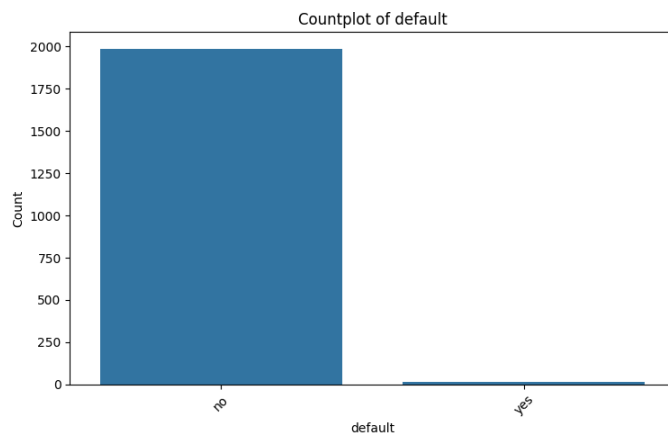


```
axes[1].set_title(f"Distribution of {column}")
```




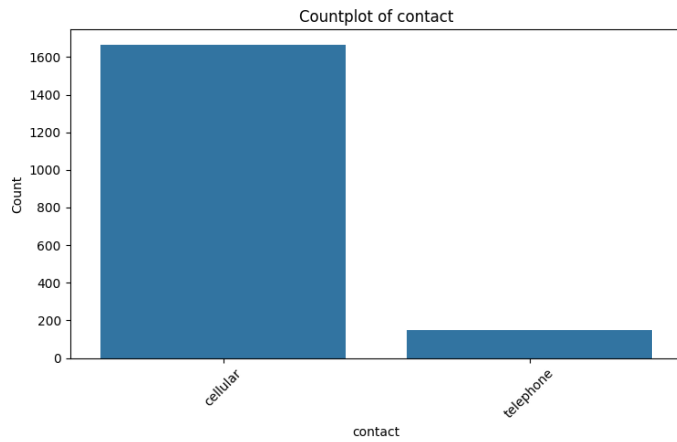


```
axes[1].set_ylabel('')
```



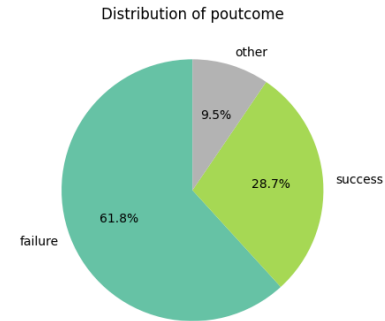
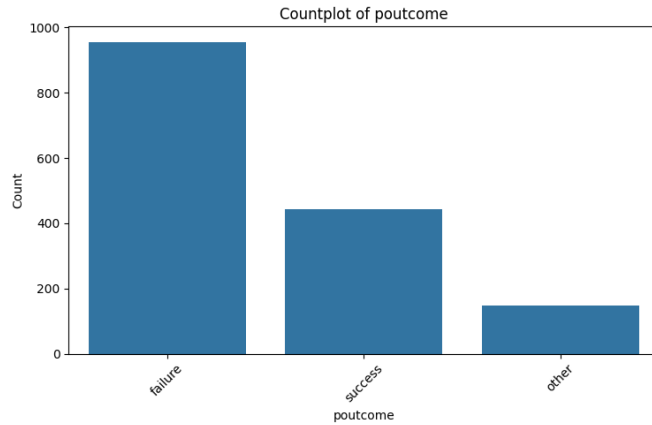
```
plt.tight_layout()
```

```
plt.show()
```









contact , default show low variaty loan show a little variaty but still might contribute to the data

1.9 targeted comparison



```
[15]: # Group by 'subscribed' and plot value counts for all categorical columns
      ↳ together
      for column in categorical_columns:
          if column != 'subscribed':
              grouped = df.groupby('subscribed')[column].value_counts(normalize=True).
              ↳ unstack().fillna(0)

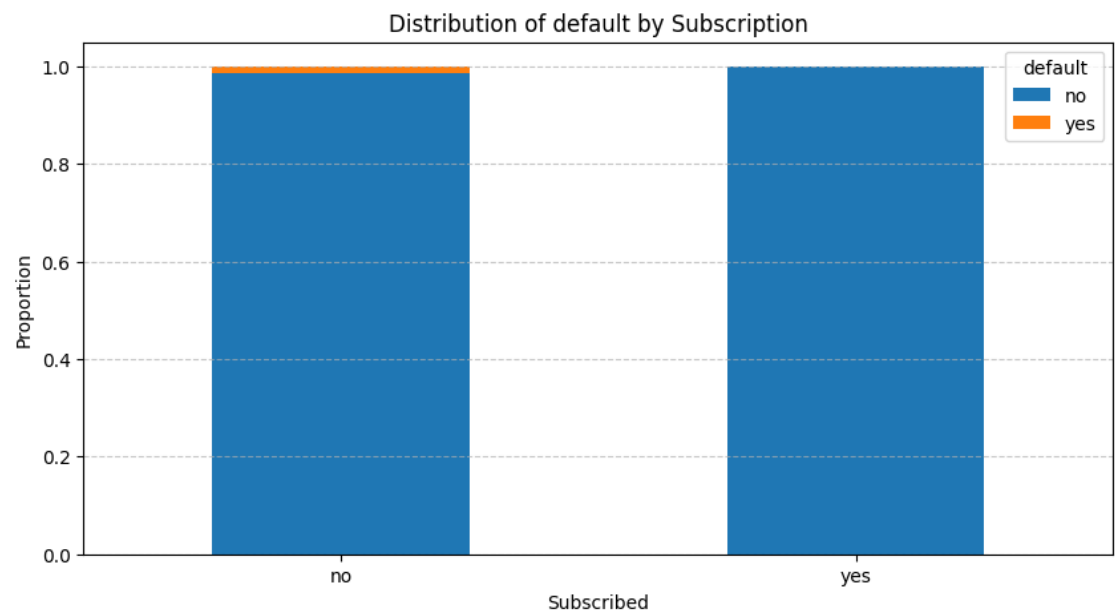
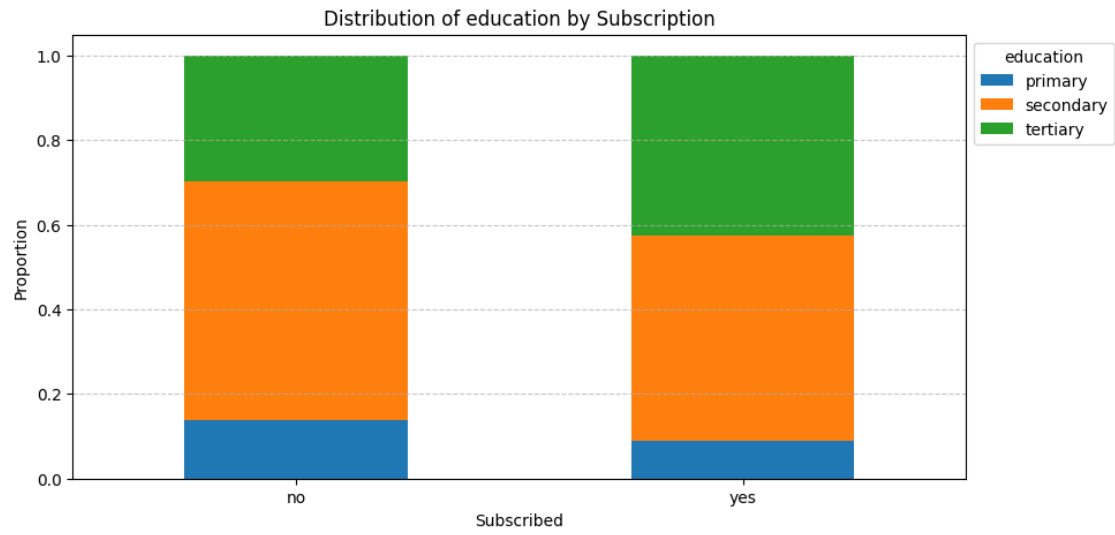
          # Plot as a stacked bar chart
          grouped.plot(kind='bar', stacked=True, figsize=(10, 5))


          plt.title(f"Distribution of {column} by Subscription")
          plt.xlabel("Subscribed")
          plt.ylabel("Proportion")
```



```
plt.legend(title=column, bbox_to_anchor=(1, 1))
plt.xticks(rotation=0)
plt.grid(axis="y", linestyle="--", alpha=0.7)


plt.show()
```











```
[16]: print("--- Numerical Columns Grouped by Subscribed ---")

for column in numerical_columns:
    print(f"\nStatistics for {column}:\n")
    display(df.groupby('subscribed')[column].describe())
```

--- Numerical Columns Grouped by Subscribed ---

Statistics for age:

| | count | mean | std | min | 25% | 50% | 75% | max |
|------------|-------|-----------|-----------|------|------|------|------|------|
| subscribed | | | | | | | | |
| no | 991.0 | 40.655903 | 9.192425 | 22.0 | 33.0 | 39.0 | 48.0 | 64.0 |
| yes | 997.0 | 42.843531 | 15.382656 | 18.0 | 31.0 | 38.0 | 54.0 | 93.0 |

Statistics for balance:

| | count | mean | std | min | 25% | 50% | 75% | max |
|------------|--------|----------|-------------|--------|--------|-------|---------|---------|
| subscribed | | | | | | | | |
| no | 1000.0 | 942.862 | 2007.134003 | -980.0 | 114.75 | 393.0 | 970.25 | 26306.0 |
| yes | 1000.0 | 1884.465 | 3891.864047 | -205.0 | 315.00 | 875.0 | 2304.50 | 81204.0 |

Statistics for day:

| | count | mean | std | min | 25% | 50% | 75% | max |
|------------|-------|------|-----|-----|-----|-----|-----|-----|
| subscribed | | | | | | | | |

| | | | | | | | | |
|-----|--------|--------|-----------|-----|-----|------|-------|------|
| no | 1000.0 | 12.364 | 10.667394 | 1.0 | 4.0 | 8.0 | 27.25 | 30.0 |
| yes | 1000.0 | 15.339 | 8.397893 | 1.0 | 9.0 | 14.0 | 22.00 | 31.0 |

Statistics for duration:

| | | | | | | | | |
|------------|--------|---------|------------|------|-------|-------|--------|--------|
| | count | mean | std | min | 25% | 50% | 75% | max |
| subscribed | | | | | | | | |
| no | 1000.0 | 206.696 | 175.152259 | 7.0 | 96.0 | 155.5 | 256.00 | 1823.0 |
| yes | 1000.0 | 377.345 | 230.154246 | 23.0 | 224.0 | 310.0 | 457.25 | 1720.0 |

Statistics for campaign:

| | | | | | | | | |
|------------|--------|-------|----------|-----|-----|-----|-----|------|
| | count | mean | std | min | 25% | 50% | 75% | max |
| subscribed | | | | | | | | |
| no | 1000.0 | 1.957 | 1.443341 | 1.0 | 1.0 | 1.0 | 2.0 | 11.0 |
| yes | 1000.0 | 1.862 | 1.310219 | 1.0 | 1.0 | 1.0 | 2.0 | 11.0 |

Statistics for pdays:

| | | | | | | | | |
|------------|--------|---------|------------|------|-------|-------|-------|-------|
| | count | mean | std | min | 25% | 50% | 75% | max |
| subscribed | | | | | | | | |
| no | 1000.0 | 185.400 | 99.759611 | -1.0 | 136.0 | 211.0 | 259.0 | 536.0 |
| yes | 1000.0 | 150.392 | 155.468012 | -1.0 | -1.0 | 123.5 | 185.0 | 854.0 |

Statistics for previous:

| | | | | | | | | |
|------------|--------|-------|----------|-----|-----|-----|-----|------|
| | count | mean | std | min | 25% | 50% | 75% | max |
| subscribed | | | | | | | | |
| no | 1000.0 | 2.362 | 3.287516 | 0.0 | 1.0 | 2.0 | 3.0 | 51.0 |
| yes | 1000.0 | 2.761 | 3.500590 | 0.0 | 0.0 | 2.0 | 4.0 | 55.0 |

1.10 targetting the pdays

```
[17]: (df['pdays'] == -1).mean()
```

```
[17]: np.float64(0.227)
```

22.7 % of the clients hasn t been contacted

1.11 targeting the balance

```
[18]: # How many zero values in `balance`?
print("There are %d account holder or %5f of the total clients who have zero_
↳balance" % ((df[df['balance']==0]['balance'].count()),

↳          (df[df['balance']==0]['balance'].count()/(df['balance'].count()))
# How many negative values in `balance`?
print("There are %d account holder or %5f of the total clients who owe money" %_
↳((df[df['balance']<0]['balance'].count()),

↳          (df[df['balance']<0]['balance'].count()/(df['balance'].count()))
```

There are 86 account holder or 0.043000 of the total clients who have zero balance

There are 93 account holder or 0.046500 of the total clients who owe money

```
[19]: # Is there any of those clients who subscribed to term deposit?
print("There are %d account holder who have zero balance and subscribed term_
↳deposit" % df[(df['balance']==0) & (df['subscribed']=='yes')]['balance'].
↳count())

print("There are %d account holder who have negative balance and subscribed_
↳term deposit" % df[(df['balance']<0) & (df['subscribed']=='yes')]['balance'].
↳count())
```

There are 42 account holder who have zero balance and subscribed term deposit

There are 7 account holder who have negative balance and subscribed term deposit

1.12 targeting duration

```
[20]: # The range of calls duration from clients who subscribed to the term deposit
print("The minimum duration (in seconds) to finalize a deal :",_
↳df[df['subscribed']=='yes']['duration'].min())
print("The maximum duration (in seconds) to finalize a deal :",_
↳df[df['subscribed']=='yes']['duration'].max())

# The average of calls duration from clients who subscribed to the term deposit
print("The average duration (in seconds) to finalize a deal :",_
↳df[df['subscribed']=='yes']['duration'].mean())
```


The minimum duration (in seconds) to finalize a deal : 23

The maximum duration (in seconds) to finalize a deal : 1720

The average duration (in seconds) to finalize a deal : 377.345

```
[21]: # The calls duration distribution from clients who subscribed to the term_
↳deposit
sns.histplot(df[df['subscribed']=='yes']['duration'], color='blue', kde=True)
```

```
[21]: <Axes: xlabel='duration', ylabel='Count'>
```



Most subscribing clients had calls lasting between 175 and 375 seconds.

!!!!!!! correlation doesn't mean causation

```
[22]: # show the clients with no contact in the last campaign but subscribed to the term deposit
      df[(df['duration']==0) & (df['subscribed']=='yes')]
```

```
[22]: Empty DataFrame
      Columns: [age, job, marital, education, default, balance, housing, loan, contact, day, month, duration, campaign, pdays, previous, poutcome, subscribed]
      Index: []
```

1.13 targeting the poutcome

```
[23]: # Show Clients who have been contacted before but unknown outcome
      df[(df['poutcome'] == '') & (df['previous']!=0)]
```

```
[23]: Empty DataFrame
      Columns: [age, job, marital, education, default, balance, housing, loan,
```

```
contact, day, month, duration, campaign, pdays, previous, poutcome, subscribed]
Index: []
```

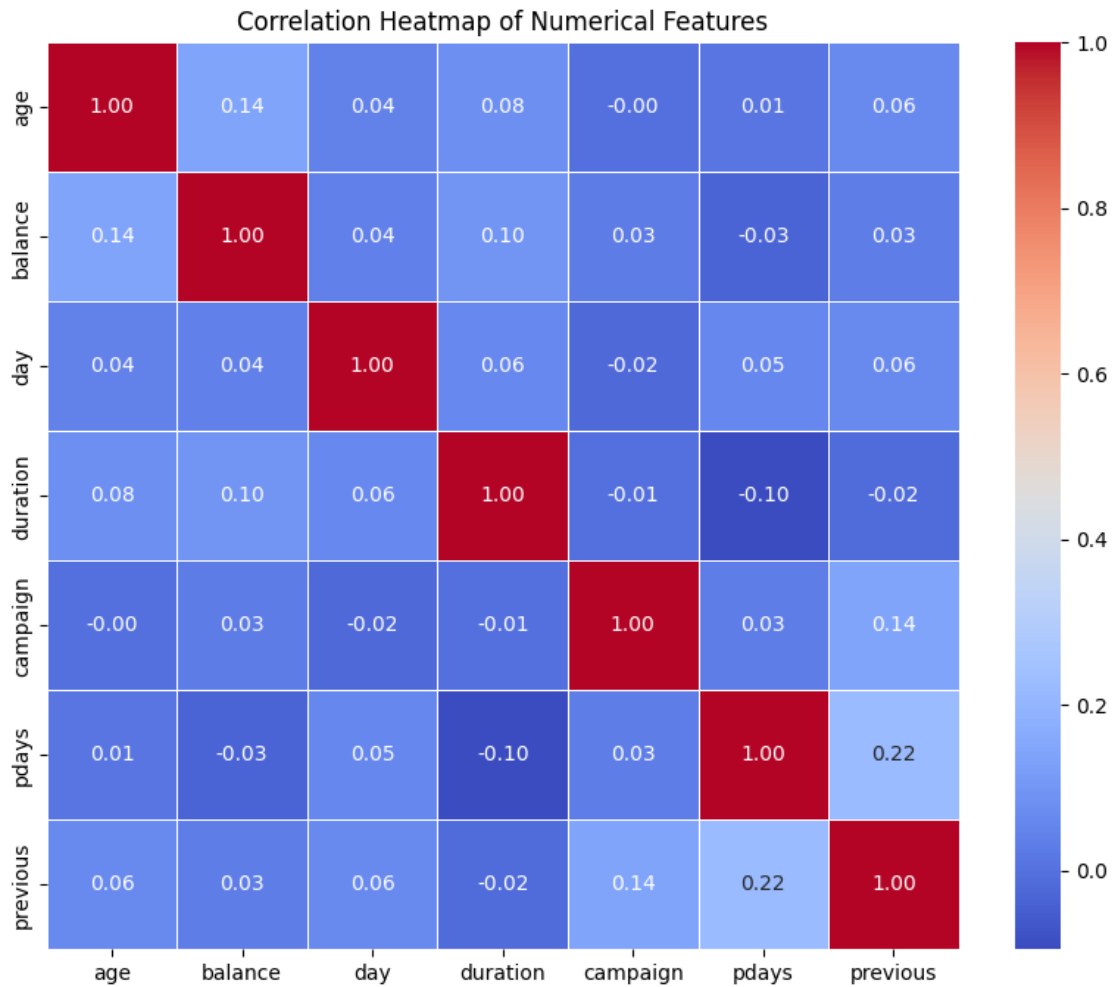
```
[24]: # pdays is -1 => previous is 0 ?
print(df[(df['pdays'] != -1) & (df['previous']==0)]['pdays'].count() ==
      df[(df['pdays'] == -1) & (df['previous'] != 0)]['pdays'].count())
```

True

1.14 correlation matrix

```
[25]: # Calculate the correlation matrix
correlation_matrix = df.select_dtypes(include=['int64', 'float64']).corr()

# Plot the heatmap
plt.figure(figsize=(10, 8))
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm', fmt='.2f',
            linewidths=0.5)
plt.title('Correlation Heatmap of Numerical Features')
plt.show()
```



there is almost no correlation between any pair of numerical columns

2 summary

nan (missing value) exists on these features: * Job * Education * Contact * Poutcome

columns to drop : * contact : Almost all clients contacted via cellular

* default : Almost no clients have defaulted (99.2% “no”)

* day : we ll treat the cycle data in month and dispose of this column * duration : has a wide range of values ,might outweigh other important features.

3 DATA preparation

```
[26]: df.isnull().sum()
```

```
[26]: age          12
      job          10
      marital      0
      education    104
      default      0
      balance      0
      housing      0
      loan         0
      contact     191
      day          0
      month        0
      duration     0
      campaign     0
      pdays       0
      previous     0
      poutcome     454
      subscribed   0
      dtype: int64
```

```
[27]: original_df = df.copy()
```

3.0.1 Dropping unnecessary columns

```
[28]: df.drop(columns=['contact', 'duration', 'default', 'day'], inplace=True)
```

```
[29]: df.columns
```

```
[29]: Index(['age', 'job', 'marital', 'education', 'balance', 'housing', 'loan',
          'month', 'campaign', 'pdays', 'previous', 'poutcome', 'subscribed'],
          dtype='object')
```

3.0.2 handling missing values

- age : the distribution is positively skewed for that we ll use the median to impute the missing values

```
[30]: df.fillna({'age': df['age'].median()}, inplace=True)
```

- poutcome : the data isn't missing but rather unknown

```
[31]: df['poutcome'] = df['poutcome'].replace("", "never")
```

- job : missing values < 0.5 % of the data , we can dispose of those rows

```
[32]: df = df.dropna(subset=['job'])
```

- education : we create a new category of unknown education level ' unknown '

```
[33]: df.fillna({'education': 'unknown'}, inplace=True)
```


3.0.3 encoding

```
[34]: # Convert 'month' into seasonal categories
season_map = {
    'dec': 'Winter', 'jan': 'Winter', 'feb': 'Winter',
    'mar': 'Spring', 'apr': 'Spring', 'may': 'Spring',
    'jun': 'Summer', 'jul': 'Summer', 'aug': 'Summer',
    'sep': 'Fall', 'oct': 'Fall', 'nov': 'Fall'
}

df['season'] = df['month'].map(season_map)
df.drop(columns=['month'], inplace=True) # Drop original month column

# One-Hot Encoding for 'season' column
df = pd.get_dummies(df, columns=['season'], drop_first=True)

# One-Hot Encoding for Other Nominal Categorical Columns
df = pd.get_dummies(df, columns=['job', 'marital', 'education', 'poutcome'],
                    drop_first=True)

[35]: # Label Encoding for Binary Categorical Columns
binary_cols = ['housing', 'loan']
le = LabelEncoder()

for col in binary_cols:
    df[col] = le.fit_transform(df[col])

[36]: max_pdays = df[df['pdays'] != -1]['pdays'].max()
df['pdays'] = df['pdays'].replace(-1, - max_pdays)
```

3.1 Detecting outliers

```
[37]: numerical_columns = [col for col in numerical_columns if col not in ['day',
    'duration']]

# Compute IQR (Interquartile Range)
Q1 = df[numerical_columns].quantile(0.25)
Q3 = df[numerical_columns].quantile(0.75)
IQR = Q3 - Q1

# Define outlier mask
outlier_mask = (df[numerical_columns] < (Q1 - 1.5 * IQR)) |
    (df[numerical_columns] > (Q3 + 1.5 * IQR))

print(outlier_mask.sum())

# Save a copy before removing outliers
```

```

df_before = df.copy()

# Remove outliers
df_after = df[~outlier_mask.any(axis=1)]

# Plot distributions before and after removing outliers
fig, axes = plt.subplots(len(numerical_columns), 2, figsize=(12, 12))

for i, col in enumerate(numerical_columns):
    # Before removing outliers
    sns.histplot(df_before[col], kde=True, ax=axes[i, 0], bins=30)
    axes[i, 0].set_title(f"{col} Distribution (Before Outlier Removal)")

    # After removing outliers
    sns.histplot(df_after[col], kde=True, ax=axes[i, 1], bins=30)
    axes[i, 1].set_title(f"{col} Distribution (After Outlier Removal)")

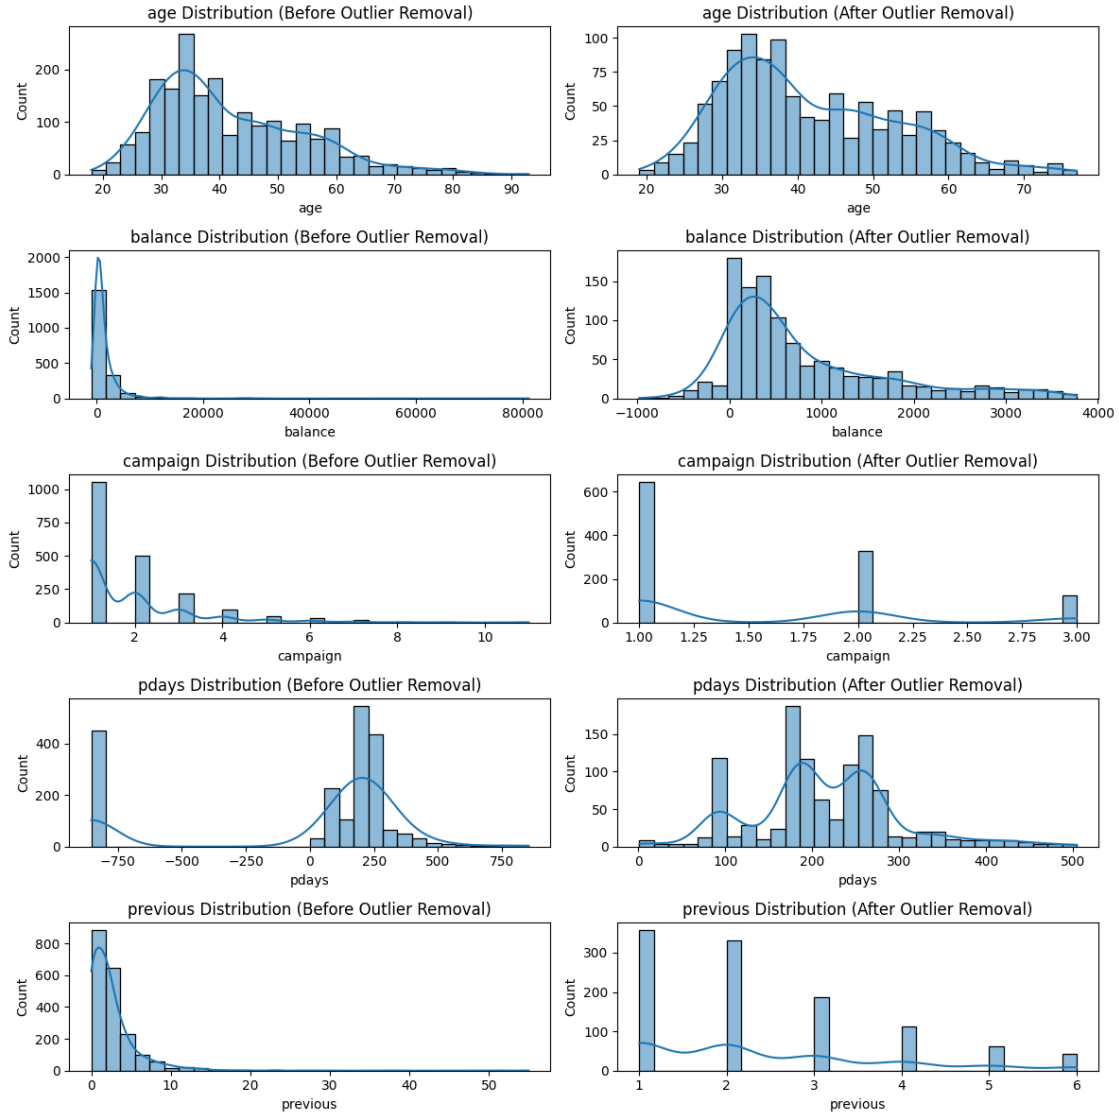
plt.tight_layout()
plt.show()

```

```

age          24
balance      159
campaign     212
pdays       483
previous     169
dtype: int64

```



3.2 Summary: Outlier Handling Decisions (Based on Distribution Analysis)

1. Features Where Outliers Were Removed We removed outliers from the following features due to extreme values affecting the distribution:

- **balance** → Previously had extreme right skew, making it difficult to analyze. Outliers removed to keep reasonable financial variations.
- **campaign** → Most clients had 1-3 contacts, but some had 10+ contacts, which were removed for better distribution.
- **previous** → Long tail distribution; extreme values (very high past contacts) were removed to avoid misleading clustering.

2. Features Where Outliers Were Kept We decided to keep outliers in these features as they hold meaningful patterns:

- **age** → The distribution remains fairly normal after outlier treatment, so extreme values were not removed.
 - **pdays** → Kept all values but handled -1 (never contacted) by replacing it with - (max value) to preserve meaning.
-

Final Decision (Based on Distribution Plots)

- **Kept & Scaled:** age, balance, campaign, pdays, previous
- **Dropped:** 'contact', 'duration', 'default', 'day'
- **Transformed pdays**

Now, we can proceed with scaling numerical features and preparing for clustering.

3.3 Scaling

```
[38]: # Define numerical columns to scale
numerical_columns = ['age', 'balance', 'campaign', 'previous', 'pdays']

# Initialize RobustScaler
scaler = RobustScaler()

# Apply scaling
df[numerical_columns] = scaler.fit_transform(df[numerical_columns])
```

3.4 dropping the target variable

```
[39]: df = df.drop(columns=['subscribed'])
```

3.5 Dimensionality Reduction

```
[46]: # Apply PCA
df.info()
pca = PCA()
pca_fit = pca.fit(df)


# Plot explained variance
plt.figure(figsize=(10, 5))
plt.plot(np.cumsum(pca.explained_variance_ratio_), marker='o', linestyle='--', color='b')
plt.xlabel('Number of Principal Components')
plt.ylabel('Cumulative Explained Variance')
plt.title('Explained Variance by Number of PCA Components')
plt.grid()
```

```
plt.show()

# Choose number of components to retain 95% variance
explained_variance = np.cumsum(pca.explained_variance_ratio_)
optimal_components = np.argmax(explained_variance >= 0.95) + 1 # First index
↳ where variance 95%

# Apply PCA with optimal components
pca_final = PCA(n_components=optimal_components)
df_pca = pca_final.fit_transform(df)
df_pca2 = df_pca.copy()
print(f"Optimal number of components: {optimal_components}")
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 1990 entries, 0 to 1999
Data columns (total 28 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   age                                    1990 non-null   float64
1   balance                               1990 non-null   float64
2   housing                               1990 non-null   int64
3   loan                                   1990 non-null   int64
4   campaign                              1990 non-null   float64
5   pdays                                1990 non-null   float64
6   previous                              1990 non-null   float64
7   season_Spring                         1990 non-null   bool
8   season_Summer                         1990 non-null   bool
9   season_Winter                         1990 non-null   bool
10  job_blue-collar                       1990 non-null   bool
11  job_entrepreneur                      1990 non-null   bool
12  job_housemaid                         1990 non-null   bool
13  job_management                        1990 non-null   bool
14  job_retired                           1990 non-null   bool
15  job_self-employed                     1990 non-null   bool
16  job_services                          1990 non-null   bool
17  job_student                           1990 non-null   bool
18  job_technician                        1990 non-null   bool
19  job_unemployed                        1990 non-null   bool
20  marital_married                       1990 non-null   bool
21  marital_single                        1990 non-null   bool
22  education_secondary                   1990 non-null   bool
23  education_tertiary                    1990 non-null   bool
24  education_unknown                     1990 non-null   bool
25  poutcome_other                         1990 non-null   bool
26  poutcome_success                       1990 non-null   bool
27  Cluster                               1990 non-null   int32
dtypes: bool(20), float64(5), int32(1), int64(2)
memory usage: 171.0 KB
```




Optimal number of components: 11

3.6 elbow method

```
[41]: # Determine the best number of clusters using the Elbow Method
wcss = []
K_range = range(1, 11) # Checking for K from 1 to 10

for k in K_range:
    kmeans = KMeans(n_clusters=k, random_state=42, n_init=10)
    kmeans.fit(df)
    wcss.append(kmeans.inertia_) # Inertia (WCSS)

# Plot the Elbow Method
plt.figure(figsize=(8, 5))
plt.plot(K_range, wcss, marker='o', linestyle='--', color='b')
plt.xlabel('Number of Clusters (K)')
plt.ylabel('Within-Cluster Sum of Squares (WCSS)')
plt.title('Elbow Method for Optimal K')
plt.grid()
plt.show()
```



```
[45]: # Set the optimal number of clusters from the Elbow Method
      optimal_k = 4 # Replace with the best k from the elbow plot

      # Apply K-Means directly to the full dataset
      kmeans = KMeans(n_clusters=optimal_k, random_state=42, n_init=100)
      df['Cluster'] = kmeans.fit_predict(df) # Assign cluster labels to original
      ↪ dataset

      # Print cluster counts
      print("Cluster distribution:\n", df['Cluster'].value_counts())

      # Get cluster centers
      cluster_centers = kmeans.cluster_centers_

      # Select first two features for visualization (since we removed PCA)
      feature_1 = df.columns[0] # First feature
      feature_2 = df.columns[1] # Second feature

      # Visualize the clusters in 2D (first two features)
      plt.figure(figsize=(8, 5))
      sns.scatterplot(x=df[feature_1], y=df[feature_2], hue=df['Cluster'],
      ↪ palette='viridis', alpha=0.7)
```

```

# Plot cluster centers with stars (*)
plt.scatter(cluster_centers[:, 0], cluster_centers[:, 1],
            marker='*', s=300, c='red', edgecolors='black', label='Cluster_
↳Centers')

# Labels and legend
plt.xlabel(feature_1)
plt.ylabel(feature_2)
plt.title(f'K-Means Clustering (k={optimal_k}) - Full Dataset')
plt.legend(title='Cluster')
plt.grid()
plt.show()


```

Cluster distribution:

```

Cluster
0      1250
1       449
3       278
2         13
Name: count, dtype: int64

```



3.7 clustering with pca

```
[47]: # Set the optimal number of clusters from the Elbow Method
      optimal_k = 3 # Replace with the best k from the elbow plot

      # Apply PCA before clustering
      pca = PCA(n_components=2) # Reduce dimensions to 2 for visualization
      df_pca = pca.fit_transform(df) # Apply PCA transformation

      # Apply K-Means on PCA-transformed data
      kmeans = KMeans(n_clusters=optimal_k, random_state=42, n_init=100)
      df['Cluster'] = kmeans.fit_predict(df_pca) # Assign cluster labels to original
      ↪dataset

      # Print cluster counts
      print("Cluster distribution:\n", df['Cluster'].value_counts())

      # Get cluster centers in PCA space
      cluster_centers = kmeans.cluster_centers_


      # Visualize clusters using PCA components
      plt.figure(figsize=(8, 5))
      sns.scatterplot(x=df_pca[:, 0], y=df_pca[:, 1], hue=df['Cluster'],
      ↪palette='viridis', alpha=0.7)

      # Plot cluster centers with stars (*)
      plt.scatter(cluster_centers[:, 0], cluster_centers[:, 1],
      ↪marker='*', s=300, c='red', edgecolors='black', label='Cluster_
      ↪Centers')

      # Labels and legend
      plt.xlabel('Principal Component 1')
      plt.ylabel('Principal Component 2')
      plt.title(f'K-Means Clustering (k={optimal_k}) - PCA Projection')
      plt.legend(title='Cluster')
      plt.grid()
      plt.show()
```

Cluster distribution:

```
Cluster
0      1528
1       449
2        13
Name: count, dtype: int64
```






3.8 overlook on our clusters

```
[ ]: # Assuming df contains the cluster labels
cluster_columns = ["age", "balance", "pdays"] # Replace with actual numerical
↳ columns

# Create subplots
fig, axes = plt.subplots(len(cluster_columns), 1, figsize=(10, 20))

for i, column in enumerate(cluster_columns):
    sns.histplot(data=df, x=column, hue="Cluster", bins=30, kde=True,
↳ palette="viridis", ax=axes[i])
    axes[i].set_title(f"Distribution of {column} by Cluster")
    axes[i].set_xlabel(column)
    axes[i].set_ylabel("Count")

plt.tight_layout()
plt.show()
```



```
[187]: # Reload dataset (assumed that df contains the cluster labels)
# You may need to reload your preprocessed dataset here if necessary






# Placeholder: Define cluster columns for visualization
cluster_columns = ["age", "balance", "campaign", "pdays", "previous"] # Adjust
↳based on available numerical columns

# Generate synthetic cluster data for visualization (Remove this if df is
↳available)
np.random.seed(42)
df = pd.DataFrame({
    "age": np.random.randint(18, 70, 300),
    "balance": np.random.randint(-500, 10000, 300),
    "campaign": np.random.randint(1, 10, 300),
    "pdays": np.random.randint(-1, 500, 300),
    "previous": np.random.randint(0, 6, 300),
    "Cluster": np.random.choice([0, 1, 2], 300) # Simulating cluster labels
})

# Create subplots
fig, axes = plt.subplots(len(cluster_columns), 1, figsize=(10, 20))

for i, column in enumerate(cluster_columns):
    sns.histplot(data=df, x=column, hue="Cluster", bins=30, kde=True,
↳palette="viridis", ax=axes[i])
    axes[i].set_title(f"Distribution of {column} by Cluster")
    axes[i].set_xlabel(column)
    axes[i].set_ylabel("Count")

plt.tight_layout()
plt.show()
```



Cluster 0 (Majority Group - Low Engagement)

This group consists mostly of clients who have lower balance amounts and minimal previous engagement with the bank. They are characterized by a lower number of previous contacts (previous) and fewer days since the last contact (pdays). They tend to have shorter campaign interactions and appear less likely to subscribe to a term deposit. This group is least likely to subscribe.

Cluster 1 (Moderate Engagement - Balanced Clients)

This group has a mix of clients with varying balance amounts, with some having significant financial assets. They show a slightly higher tendency to have been contacted multiple times before. Their pdays values indicate that they have been previously engaged, but not necessarily in a consistent manner. This group has a moderate probability of subscribing.

Cluster 2 (Highly Engaged Clients - Most Likely to Subscribe)

Clients in this cluster generally have higher balances. They have been contacted before and show a strong engagement pattern in terms of the number of previous contacts (previous). Many of them have high pdays values, indicating that they were engaged in previous campaigns before being contacted again. They tend to have longer call durations and more interactions with the bank. This group is the most likely to subscribe to the term deposit after the campaign.

```
[48]: # Set the optimal number of clusters from the Elbow Method
      optimal_k = 2 # Replace with the best k from the elbow plot

      # Apply K-Means
      kmeans = KMeans(n_clusters=optimal_k, random_state=42, n_init=100)
      df['Cluster'] = kmeans.fit_predict(df_pca) # Assign cluster labels to original
      ↪dataset

      # Print cluster counts
      print(df['Cluster'].value_counts())

      # Get cluster centers
      cluster_centers = kmeans.cluster_centers_

      # Visualize the clusters using PCA (2D plot)
      plt.figure(figsize=(8, 5))
      sns.scatterplot(x=df_pca[:, 0], y=df_pca[:, 1], hue=df['Cluster'],
      ↪palette='viridis', alpha=0.7)

      # Plot cluster centers with stars (*)
      plt.scatter(cluster_centers[:, 0], cluster_centers[:, 1],
      ↪marker='*', s=300, c='red', edgecolors='black', label='Cluster
      ↪Centers')

      # Labels and legend
```


```
plt.xlabel('Principal Component 1')
plt.ylabel('Principal Component 2')
plt.title(f'K-Means Clustering (k={optimal_k}) - PCA Projection')
plt.legend(title='Cluster')
plt.grid()
plt.show()
```

Cluster

0 1539

1 451

Name: count, dtype: int64






```
[49]: # Assuming df contains the cluster labels
cluster_columns = ["age", "balance", "pdays"] # Replace with actual numerical columns

# Create subplots
fig, axes = plt.subplots(len(cluster_columns), 1, figsize=(10, 20))

for i, column in enumerate(cluster_columns):
    sns.histplot(data=df, x=column, hue="Cluster", bins=30, kde=True, palette="viridis", ax=axes[i])
    axes[i].set_title(f"Distribution of {column} by Cluster")
```

```
axes[i].set_xlabel(column)
axes[i].set_ylabel("Count")

plt.tight_layout()
plt.show()
```

```

[51]: # Reload dataset (assumed that df contains the cluster labels)
# You may need to reload your preprocessed dataset here if necessary

# Placeholder: Define cluster columns for visualization
cluster_columns = ["age", "balance", "campaign", "pdays", "previous"] # Adjust
↳ based on available numerical columns






# Generate synthetic cluster data for visualization (Remove this if df is
↳ available)
np.random.seed(42)
df = pd.DataFrame({
    "age": np.random.randint(18, 70, 300),
    "balance": np.random.randint(-500, 10000, 300),
    "campaign": np.random.randint(1, 10, 300),
    "pdays": np.random.randint(-1, 500, 300),
    "previous": np.random.randint(0, 6, 300),
    "Cluster": np.random.choice([0, 1], 300) # Simulating cluster labels
})

# Create subplots
fig, axes = plt.subplots(len(cluster_columns), 1, figsize=(10, 20))

for i, column in enumerate(cluster_columns):
    sns.histplot(data=df, x=column, hue="Cluster", bins=30, kde=True,
↳ palette="viridis", ax=axes[i])
    axes[i].set_title(f"Distribution of {column} by Cluster")
    axes[i].set_xlabel(column)
    axes[i].set_ylabel("Count")

plt.tight_layout()
plt.show()

```



Customers with a higher balance

The cluster with higher account balances is more likely to subscribe. These individuals have better financial stability and are more open to long-term investments like term deposits. Customers who have been contacted before

The cluster where individuals had previous successful contacts with the bank has a higher likelihood of subscribing. This suggests that prior engagement with the bank builds trust and increases the chances of conversion. Customers with a lower number of contacts per campaign

The cluster with fewer campaign contacts but a higher conversion rate is more promising. This indicates that those who subscribe often require fewer interactions, meaning they were already interested or convinced early in the campaign. Customers with medium age range (not too young, not too old)

The middle-aged group (30-50 years old) is more likely to subscribe compared to younger individuals. This demographic tends to be in a financially stable position, planning for savings and investments.

Cluster 0: The Less Engaged Group

Generally consists of clients with lower balance and less frequent previous contacts. Campaign contact frequency is lower, meaning they haven't been reached out as much. Many in this cluster have high pdays values, meaning they have not been contacted recently. This cluster represents clients who are less likely to subscribe to a term deposit because they have low interaction history with the bank and less engagement in past marketing campaigns.

Cluster 1: The More Engaged Group

Clients in this cluster tend to have a higher account balance and more frequent past contacts. The campaign contact frequency is higher, meaning they have been targeted multiple times. Their pdays values are lower, indicating they were contacted more recently. This cluster is more likely to subscribe to the term deposit since they have higher engagement with the bank and have been in recent contact.